

Deep Learning

04 Gradient Descent

Dr. Konda Reddy Mopuri
Dept. of Artificial Intelligence
IIT Hyderabad
Jan-May 2023

- ① Learning: finding a good function f^* from a set of functions \mathcal{F}

- ① Learning: finding a good function f^* from a set of functions \mathcal{F}
- ② How to find the goodness of a function f ?

- ① Learning: finding a good function f^* from a set of functions \mathcal{F}
- ② How to find the goodness of a function f ?
- ③ Through a loss $l : \mathcal{F} \times \mathcal{L} \rightarrow \mathcal{R}$

- ① Learning: finding a good function f^* from a set of functions \mathcal{F}
- ② How to find the goodness of a function f ?
- ③ Through a loss $l : \mathcal{F} \times \mathcal{L} \rightarrow \mathcal{R}$
- ④ Such that value of $l(f, z)$ increases with the *wrongness* of f on z :
(measure of discrepancy between the expected and predicted)

- ① Learning: finding a good function f^* from a set of functions \mathcal{F}
- ② How to find the goodness of a function f ?
- ③ Through a loss $l : \mathcal{F} \times \mathcal{L} \rightarrow \mathcal{R}$
- ④ Such that value of $l(f, z)$ increases with the *wrongness* of f on z :
(measure of discrepancy between the expected and predicted)
- ⑤
 - Regression: $l(f, (x, y)) = (f(x) - y)^2$
 - Classification: $l(f, (x, y)) = \mathbf{1}(f(x) \neq y)$
 - Density estimation: $l(q, z) = -\log(q(z))$

- ① Learning: finding a good function f^* from a set of functions \mathcal{F}
- ② How to find the goodness of a function f ?
- ③ Through a loss $l : \mathcal{F} \times \mathcal{L} \rightarrow \mathcal{R}$
- ④ Such that value of $l(f, z)$ increases with the *wrongness* of f on z :
(measure of discrepancy between the expected and predicted)
- ⑤
 - Regression: $l(f, (x, y)) = (f(x) - y)^2$
 - Classification: $l(f, (x, y)) = \mathbf{1}(f(x) \neq y)$
 - Density estimation: $l(q, z) = -\log(q(z))$
- ⑥ Loss may have additional terms (from prior knowledge)

Expected Risk

- ① We want f with small *expected (average) risk* $R(f) = \mathbb{E}_z(l(f, z))$

Expected Risk

- ① We want f with small *expected (average) risk* $R(f) = \mathbb{E}_z(l(f, z))$
- ② $f^* = \operatorname{argmin}_{f \in \mathcal{F}} R(f)$

Expected Risk

- ① We want f with small *expected (average) risk* $R(f) = \mathbb{E}_z(l(f, z))$
- ② $f^* = \underset{f \in \mathcal{F}}{\operatorname{argmin}} R(f)$
- ③ This is unknown. However, if the training data $\mathcal{D} = \{z_1, \dots, z_N\}$ is i.i.d. we can estimate the risk empirically (known as empirical risk),

$$\hat{R}(f; \mathcal{D}) = \hat{\mathbb{E}}_{\mathcal{D}}(l(f, z)) = \frac{1}{N} \sum_{i=1}^N l(f, z_n)$$

- How to find the model parameters that minimize the loss function?

$$w^* = \underset{w}{\operatorname{argmin}} L(w)$$

- How to find the model parameters that minimize the loss function?

$$w^* = \underset{w}{\operatorname{argmin}} L(w)$$

- General and vast, but we will discuss within our context

- Finding the parameters that minimize the training loss

$$W^*, \mathbf{b}^* = \underset{W, \mathbf{b}}{\operatorname{argmin}} \mathcal{L}(f(\cdot; W, \mathbf{b}); \mathcal{D})$$

- Finding the parameters that minimize the training loss

$$W^*, \mathbf{b}^* = \underset{W, \mathbf{b}}{\operatorname{argmin}} \mathcal{L}(f(\cdot; W, \mathbf{b}); \mathcal{D})$$

- How do we find these optimal parameters?

- Finding the parameters that minimize the training loss

$$W^*, \mathbf{b}^* = \underset{W, \mathbf{b}}{\operatorname{argmin}} \mathcal{L}(f(\cdot; W, \mathbf{b}); \mathcal{D})$$

- How do we find these optimal parameters?
 - Closed form solution (e.g. linear regression)

- Finding the parameters that minimize the training loss

$$W^*, \mathbf{b}^* = \underset{W, \mathbf{b}}{\operatorname{argmin}} \mathcal{L}(f(\cdot; W, \mathbf{b}); \mathcal{D})$$

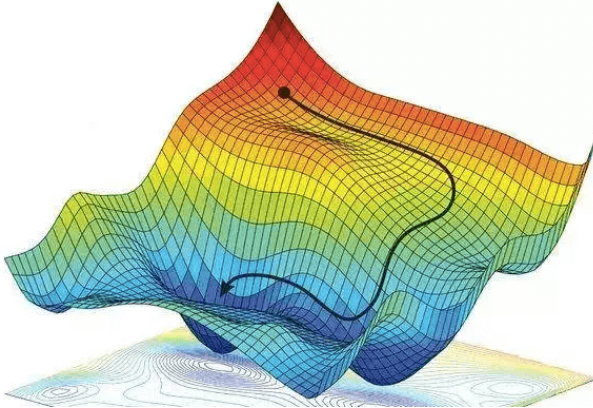
- How do we find these optimal parameters?
 - Closed form solution (e.g. linear regression)
 - Ad-hoc recipes (e.g. Perceptron, K-NN classifier)

- Finding the parameters that minimize the training loss

$$W^*, \mathbf{b}^* = \operatorname{argmin}_{W, \mathbf{b}} \mathcal{L}(f(\cdot; W, \mathbf{b}); \mathcal{D})$$

- How do we find these optimal parameters?
 - Closed form solution (e.g. linear regression)
 - Ad-hoc recipes (e.g. Perceptron, K-NN classifier)
 - What if the loss function can't be minimized analytically?

Loss surface



Source: Medium

Not-so-intelligent idea!

- Probe random directions

Not-so-intelligent idea!

- Probe random directions
- Progress if you find a useful direction

Not-so-intelligent idea!

- Probe random directions
- Progress if you find a useful direction
- Repeat

Not-so-intelligent idea!

- Probe random directions
- Progress if you find a useful direction
- Repeat
- **Very ineffective!**

A better looking one: Follow the slope!



భారతీయ టెక్నాలజీ హైదరాబాద్
भारतीय प्रौद्योगिकी संस्थान हैदराबाद
Indian Institute of Technology Hyderabad

- Sense the slope around the feet

A better looking one: Follow the slope!

- Sense the slope around the feet
- Identify the steepest direction, make a brief progress

A better looking one: Follow the slope!

- Sense the slope around the feet
- Identify the steepest direction, make a brief progress
- Repeat until convergence

A better looking one: Follow the slope!

- Sense the slope around the feet
- Identify the steepest direction, make a brief progress
- Repeat until convergence
- This is Gradient Descent!

Derivative and Gradient

- In 1D, derivative of a function gives the slope

$$\frac{\partial f}{\partial x} = \lim_{\delta \rightarrow 0} \frac{f(x + \delta) - f(x)}{h}$$

Derivative and Gradient

- In 1D, derivative of a function gives the slope

$$\frac{\partial f}{\partial x} = \lim_{\delta \rightarrow 0} \frac{f(x + \delta) - f(x)}{h}$$

- In higher dimensions, given a function

$$f : \mathcal{R}^D \rightarrow \mathcal{R}$$

gradient is the mapping

$$\begin{aligned} \nabla f : \mathcal{R}^D &\rightarrow \mathcal{R}^D \\ x &\rightarrow \left(\frac{\partial f}{\partial x_1}, \dots, \frac{\partial f}{\partial x_D} \right) \end{aligned}$$

Derivative and Gradient

- In 1D, derivative of a function gives the slope

$$\frac{\partial f}{\partial x} = \lim_{\delta \rightarrow 0} \frac{f(x + \delta) - f(x)}{h}$$

- In higher dimensions, given a function

$$f : \mathcal{R}^D \rightarrow \mathcal{R}$$

gradient is the mapping

$$\begin{aligned} \nabla f : \mathcal{R}^D &\rightarrow \mathcal{R}^D \\ x &\rightarrow \left(\frac{\partial f}{\partial x_1}, \dots, \frac{\partial f}{\partial x_D} \right) \end{aligned}$$

- ∇f vector gives the direction and rate of fastest increase for f .

$$\begin{aligned}\mathcal{L}(w + \eta u) &= \mathcal{L}(w) + \eta u^T \nabla_w \mathcal{L}(w) + \frac{\eta^2}{2!} u^T \nabla^2 \mathcal{L}(w) u + \dots \\ &\approx \mathcal{L}(w) + \eta u^T \nabla_w \mathcal{L}(w)\end{aligned}$$

- For $\mathcal{L}(w + \eta u)$ to be lesser than $\mathcal{L}(w)$, $u^T \nabla_w \mathcal{L}(w) < 0$

$$\begin{aligned}\mathcal{L}(w + \eta u) &= \mathcal{L}(w) + \eta u^T \nabla_w \mathcal{L}(w) + \frac{\eta^2}{2!} u^T \nabla^2 \mathcal{L}(w) u + \dots \\ &\approx \mathcal{L}(w) + \eta u^T \nabla_w \mathcal{L}(w)\end{aligned}$$

- For $\mathcal{L}(w + \eta u)$ to be lesser than $\mathcal{L}(w)$, $u^T \nabla_w \mathcal{L}(w) < 0$
- The difference would be least if u is in the opposite direction to $\nabla_w \mathcal{L}(w)$, the gradient

Gradient Descent

- Goal is to minimize the error (or loss): determine the parameters w that minimize the loss $\mathcal{L}(w)$

Gradient Descent

- Goal is to minimize the error (or loss): determine the parameters w that minimize the loss $\mathcal{L}(w)$
- Gradient points uphill \rightarrow negative of gradient points downhill

Gradient Descent




Figure credits: Ahmed Fawzy Gad

Gradient Descent

- ① Start with an arbitrary initial parameter vector w_0

Gradient Descent

- ① Start with an arbitrary initial parameter vector w_0
- ② Repeatedly modify it via updating in small steps

Gradient Descent

- ① Start with an arbitrary initial parameter vector w_0
- ② Repeatedly modify it via updating in small steps
- ③ At each step, modify in the direction that produces steepest descent along the error surface

How to compute the gradient?

- Numerically, for each component of w using the derivative formula

$$\frac{\partial f}{\partial x} = \lim_{\delta \rightarrow 0} \frac{f(x + \delta) - f(x)}{\delta}$$

How to compute the gradient?

- Numerically, for each component of w using the derivative formula

$$\frac{\partial f}{\partial x} = \lim_{\delta \rightarrow 0} \frac{f(x + \delta) - f(x)}{\delta}$$

- Slow and approximate!

How to compute the gradient?

- Analytically, using calculus for computing the derivatives

$$L_i = \sum_{j \neq y_i} \max\{0, s_j - s_{y_i} + 1\}$$

$$L = \frac{1}{N} \sum_i L_i + \sum_k w_k^2$$

$$s = f(x, W)$$

$$\nabla L_{iw}?$$

How to compute the gradient?

- Analytically, using calculus for computing the derivatives

$$L_i = \sum_{j \neq y_i} \max\{0, s_j - s_{y_i} + 1\}$$

$$L = \frac{1}{N} \sum_i L_i + \sum_k w_k^2$$

$$s = f(x, W)$$

$$\nabla L_{iw}?$$

- Analytic way is fast, exact, but error-prone!

Batch Gradient Descent

```
for i in range(nb_epochs):  
     $\nabla L_w$  = evaluate_gradient(L,  $\mathcal{D}$ , w)  
    w = w -  $\eta$  *  $\nabla L_w$ 
```

Batch Gradient Descent

```
for i in range(nb_epochs):  
     $\nabla L_w$  = evaluate_gradient(L,  $\mathcal{D}$ , w)  
    w = w -  $\eta$  *  $\nabla L_w$ 
```

- ① Guaranteed to converge to global minima in case of convex functions, and to a local minima in case of non-convex functions

Stochastic Gradient Descent (SGD)

- ① Performs updates parameters for each training example

$$w = w - \eta \nabla_w \mathcal{L}(w, x^i, y^i)$$

Stochastic Gradient Descent (SGD)

- ① Performs updates parameters for each training example
$$w = w - \eta \nabla_w \mathcal{L}(w, x^i, y^i)$$
- ② In case of large datasets, Batch GD computes redundant gradients for similar examples for each parameter update

Stochastic Gradient Descent (SGD)

- ① Performs updates parameters for each training example
$$w = w - \eta \nabla_w \mathcal{L}(w, x^i, y^i)$$
- ② In case of large datasets, Batch GD computes redundant gradients for similar examples for each parameter update
- ③ SGD does away with redundancy and generally faster and can be used to learn online

Stochastic Gradient Descent (SGD)

- ① However, frequent updates with a high variance cause the objective function to fluctuate heavily




Figure credits: Wikipedia

Stochastic Gradient Descent (SGD)

- ① SGD's fluctuations enable it to jump to new and potentially better local minima

Stochastic Gradient Descent (SGD)

- ① SGD's fluctuations enable it to jump to new and potentially better local minima
- ② This complicates the convergence, as it overshoots

Stochastic Gradient Descent (SGD)

- ① SGD's fluctuations enable it to jump to new and potentially better local minima
- ② This complicates the convergence, as it overshoots
- ③ However, if the learning rate is slowly decreased, we can show similar convergence to Batch GD

Stochastic Gradient Descent (SGD)

```
for i in range(nb_epochs):  
    np.random.shuffle( $\mathcal{D}$ )  
    for  $x_i \in \mathcal{D}$ :  
         $\nabla L_w = \text{evaluate\_gradient}(L, x_i, w)$   
         $w = w - \eta * \nabla L_w$ 
```

Mini-batch Gradient Descent

- ① Takes the best of both worlds, updates the parameters for every mini-batch of n samples

$$w = w - \eta \nabla_w \mathcal{L}(w, x^{i:i+n}, y^{i:i+n})$$

Mini-batch Gradient Descent

- ① Takes the best of both worlds, updates the parameters for every mini-batch of n samples

$$w = w - \eta \nabla_w \mathcal{L}(w, x^{i:i+n}, y^{i:i+n})$$

- ②
 - Reduces the variance of the parameter updates, which can lead to more stable convergence
 - Can make use of highly optimized matrix optimizations

Mini-batch Gradient Descent

- ① Takes the best of both worlds, updates the parameters for every mini-batch of n samples
$$w = w - \eta \nabla_w \mathcal{L}(w, x^{i:i+n}, y^{i:i+n})$$
- ②
 - Reduces the variance of the parameter updates, which can lead to more stable convergence
 - Can make use of highly optimized matrix optimizations
- ③ Common mini-batch sizes vary from 32 to 1024, depending on the application

Mini-batch Gradient Descent

- ① Takes the best of both worlds, updates the parameters for every mini-batch of n samples
$$w = w - \eta \nabla_w \mathcal{L}(w, x^{i:i+n}, y^{i:i+n})$$
- ②
 - Reduces the variance of the parameter updates, which can lead to more stable convergence
 - Can make use of highly optimized matrix optimizations
- ③ Common mini-batch sizes vary from 32 to 1024, depending on the application
- ④ This is the algorithm of choice while training DNNs (also, incorrectly referred to as SGD in general)

Mini-batch Gradient Descent

```
for i in range(nb_epochs):  
    np.random.shuffle( $\mathcal{D}$ )  
    for batch in get_batches( $\mathcal{D}$ , batch_size = 128):  
         $\nabla L_w$  = evaluate_gradient(L, batch, w)  
         $w = w - \eta * \nabla L_w$ 
```


Some challenges

- ① Choosing a proper learning rate

Some challenges

- ① Choosing a proper learning rate
 - Learning rate schedules try to adjust it during the training

Some challenges

- ① Choosing a proper learning rate
 - Learning rate schedules try to adjust it during the training
 - However, these schedules are defined in advance and hence unable to adapt to the task at hand

Some challenges

- ① Choosing a proper learning rate
 - Learning rate schedules try to adjust it during the training
 - However, these schedules are defined in advance and hence unable to adapt to the task at hand
- ② Same learning rate applies to all the parameters

Some challenges

- ① Choosing a proper learning rate
 - Learning rate schedules try to adjust it during the training
 - However, these schedules are defined in advance and hence unable to adapt to the task at hand
- ② Same learning rate applies to all the parameters
- ③ Avoiding numerous sub-optimal local minima